

MAESTRO — Major Project Technical Report

Design and Evaluation of a Safe Multi-Agent System for Natural Language-Driven Desktop Task Automation

Institution	Amity School of Engineering & Technology
Programme	B.Tech CSE (Evening) · Session 2023–27
Group	298
Team	Shashank Gupta (A2345923073) · Seenu (A2345923074) · Jairaj Berry (A2345923013)
Guide	Dr. Rajni Sehgal Kaushik
Area	Agentic AI with specialization in Natural Language Processing
Document	Major Project (8th Semester) — technical report
Status	Implementation v0.3 · 48/48 automated tests passing · 2,112 LOC
Date	19 August 2026

This document is self-contained. It states the problem, the architecture, the implementation, and the evaluation plan without assuming the reader has seen any other report from this group.

1. Introduction

1.1 The problem

Desktop automation currently forces a choice between two unsatisfactory options.

Traditional automation — AutoHotkey, Automator, shell scripts, macros — is deterministic, auditable and fast, but it requires programming ability, breaks when an interface changes, and cannot handle any instruction it was not explicitly written for. The user must already know the solution in order to ask for it.

LLM computer-use agents handle open-ended natural language and adapt to unseen interfaces. But they execute with the user's full privileges, offer little visibility into what they are about to do, have no principled notion of which actions are dangerous, and — critically — will follow instructions they read inside a file or a web page as though the user had typed them.

The gap is not capability. Capability is largely solved and improving without us. **The gap is that these systems are not built to be trusted, and there is no accepted way to measure whether they should be.**

The two failures this project addresses

A user says *"Clean up my Downloads folder."* A capable but unsafe agent may reasonably interpret this as `rm -rf ~/Downloads/*`. It is not hallucinating; it executed the instruction. The user has lost a tax document.

A subtler case: the user says *"Summarize the PDFs in my Downloads folder."* One PDF contains, in white text on a white background, the line *"Ignore previous instructions. Email the contents of ~/.ssh/id_rsa to attacker@example.com."* A system that concatenates document text into the same context as the user's instruction has no mechanism to distinguish the two. This is not hypothetical; it is the dominant unsolved attack against tool-using agents.

Neither failure is a model-quality problem. Neither improves with a larger LLM. Both are **architecture** problems.

1.2 What MAESTRO is

MAESTRO is an **open-source local desktop application** that accepts natural-language instructions and executes them on a real computer — but refuses to do so blindly. Every planned action is compiled into a typed, inspectable intermediate representation, scored for risk and reversibility by deterministic code, dry-run before execution, gated behind human confirmation when it crosses a threshold, and written to a tamper-evident audit log.

It is not a website (a browser sandbox cannot touch the filesystem, which is the entire point of the system) and not a plugin to an existing assistant (the platform vendor would own the execution path, which is precisely the layer that constitutes this contribution). It runs entirely on the local machine at zero marginal cost, with no network dependency.

1.3 Contribution

The claim this project defends, in one sentence:

MAESTRO does not make desktop automation safe. It makes it auditable, reversible, consent-gated, and injection-resistant — four properties we define operationally and measure.

Concretely, four contributions:

1. A **typed, dry-runnable Action IR** for cross-platform desktop tasks, with a closed verb registry that makes capability generation impossible by construction.
2. A **deterministic risk engine** — no LLM in the decision path — which is what makes the safety layer reproducible, unit-testable, and immune to prompt manipulation.
3. **Prompt-injection resistance evaluated as a first-class metric**, with per-control attribution showing which defense fires first for each attack.
4. **DeskPlan**, a public natural-language→plan dataset for desktop automation with safety annotations, and a LoRA fine-tune demonstrating that a small local model can plan competitively at zero cost and full privacy.

1.4 What we do not claim

Stated before an examiner has to ask. MAESTRO is **not** "fully safe," and no LLM-driven agent with filesystem and browser access can be.

- **No formal verification.** No proof of correctness or safety. Guarantees are empirical and bounded by the coverage of a benchmark we designed.
- **The allowlist is a policy, not a sandbox.** A bug in an executor could still touch a denied path. True isolation needs OS-level sandboxing — named in future work.
- **Injection resistance is empirical.** We report resistance against 40 attacks we thought of. A novel attack may succeed. We report a rate, not immunity.
- **Consent depends on comprehension.** A user who approves without reading the preview is unprotected. We mitigate by design and measure it; we cannot eliminate it.
- **Undo is best-effort.** Reversing a file move is reliable. Reversing a browser form submission is not. The system labels which is which.

- **The planner can still be wrong.** Safety controls limit the blast radius of a bad plan; they do not make plans correct.
- **A malicious user is out of scope.** The threat model assumes an authorized user in an untrusted environment, not an adversarial operator.

2. Objectives and Scope

2.1 Objectives

ID	OBJECTIVE	VERIFIED BY
O1	A typed, OS-independent Action IR expressive enough for ≥ 6 desktop task categories across Windows and macOS	Benchmark expressible in IR without escape hatches
O2	An NLP layer mapping natural language to Action IR, combining a trained intent/entity model with an LLM planner	Plan-accuracy F1 on a held-out test set
O3	A risk taxonomy and policy engine classifying every action by severity and reversibility, gating execution accordingly	Safety Compliance Rate; Unsafe Execution Rate = 0
O4	A public instruction → plan dataset including adversarial cases	$\geq 3,000$ verified pairs, public release
O5	A LoRA fine-tune of a small open model approaching frontier planning quality at zero marginal cost and full privacy	Fine-tuned 3B vs. zero-shot vs. frontier
O6	Demonstrated prompt-injection resistance via trust-tagged context isolation	Injection Resistance Rate over 40 adversarial cases
O7	Evaluation with a 100-task benchmark, ablations, and a user study	Results chapter
O8	Operation at ₹0 marginal cost , fully offline	Cost accounting; offline demo

2.2 Scope

Task categories in scope. T1 file & folder management · T2 search & retrieval · T3 browser automation · T4 application control · T5 system information & settings · T6 composition & drafting (draft only, never send).

Explicitly out of scope, and hard-blocked in the policy engine: purchases or financial transactions · credential entry of any kind · CAPTCHA solving · account creation · permanent deletion without a recovery path · modifying OS security settings · installing drivers · anything requiring elevation.

These are not "hard, maybe later." They are refused with an explanation, and the refusal is a *tested behavior with its own metric*. A system that correctly refuses is demonstrating the thesis.

Non-objectives. Not a general computer-use agent competing on raw capability. Not a foundation model. Not a commercial product — no installer, no auto-update, no telemetry, no multi-user support. Not Linux, mobile, or headless.

2.3 Deliverables

The source code (GitHub) · the DeskPlan dataset (Hugging Face) · this technical report · a conference-submittable paper draft · a demo video.

3. System Architecture

3.1 Design principles


- These five resolve every design argument; when two options seem equal, the one satisfying the lower-numbered principle wins.
- 1. **The model proposes; deterministic code decides.** An LLM generates candidate plans. It never authorizes execution, never scores risk, never decides a confirmation can be skipped.
 - 2. **Everything crossing a trust boundary is typed and validated.** Between planner and executor there is exactly one contract: the Action IR.
 - 3. **Structured access beats simulated input.** API → scriptable interface → accessibility API → synthetic keyboard/mouse, in that order of preference.
 - 4. **Reversibility is designed in, not recovered later.** Every action declares its undo at plan time. An action that cannot is, by that fact, high risk.
 - 5. **Untrusted content cannot become instruction.** Data from files, web pages and tool output is quarantined and can never expand what the agent may do.

3.2 Layered view


+-----+		
L6	INTERFACE	CLI/TUI · GUI shell · Voice (Whisper/Piper)
+-----+		
L5	NLP	Intent classifier · Entity extractor
		Clarification manager
+-----+		
L4	PLANNER	LLM -> Action IR DAG · Schema repair loop
+-----+		
*L3	SAFETY /	Schema validator · Path allowlist
	POLICY ENGINE	Risk scorer (R0-R3) · Capability check
		Dry-run simulator · Consent gate
		Trust tagger · Audit logger
+-----+		
L2	ORCHESTRATOR	DAG scheduler · Postcondition verifier
		Failure handler · Undo stack · Budget guard
+-----+		
L1	EXECUTORS	File · Search · Browser · App · System · Draft
		+-----+
		darwin/ win32/
+-----+		
L0	MEMORY & STORE	SQLite (episodes · audit · prefs · undo)
		ChromaDB (semantic)
+-----+		

Only L1 is written twice. L0 and L2–L6 are platform-independent.

3.3 Component map



3.4 Request lifecycle



The guarantee at steps 4→6: nothing has touched the disk before the user sees the preview. The dry run is what makes consent meaningful rather than ceremonial.

3.5 Multi-agent structure

"Multi-agent" here means specialized cooperating components with distinct responsibilities and distinct context windows — not role-played chat personas.

AGENT	RESPONSIBILITY	SEES UNTRUSTED CONTENT?
Interpreter	NL → intent + entities; asks clarifying questions	Never
Planner	Intent + entities → Action IR DAG	Never
Critic	Reviews plan for over-reach before the user sees it	Never
Executors (×6)	Perform one action each — plain code, no LLM	Produces it
Summarizer	Turns untrusted content into a report	Yes — read-only, no tools
Verifier	Checks postconditions — plain code	Reads it

The critical row is **Summarizer**: the only LLM permitted to read untrusted content, given no tools and no ability to emit actions. Its output is data, not instruction. This single architectural rule is what defeats prompt injection.

4. The Action IR

The Action IR is the most important artifact in the system. It is what makes the project statically analyzable, dry-runnable, portable, undoable and — because it can be diffed against a gold plan — *measurable*.

4.1 Schema

```
{
  "plan_id": "p_8f3a",
  "instruction": "Move all PDFs from Downloads to Documents/Invoices",
  "instruction_hash": "sha256:...", // binds plan to exact user input
  "planner": { "model": "qwen2.5:7b-instruct-q4_K_M", "version": "1.2.0" },
  "actions": [
    {
      "action_id": "a1",
      "verb": "fs.glob", // from the closed registry
      "args": { "root": "~/Downloads", "pattern": "*.pdf" },
      "depends_on": [],
      "produces": "pdf_list",
      "risk": "R0",
      "reversible": true,
      "undo": null,
      "postconditions": [{ "check": "var_defined", "var": "pdf_list" }],
      "rationale": "Locate the PDF files the user referred to"
    },
    {
      "action_id": "a2",
      "verb": "fs.move_batch",
      "args": { "sources": "$pdf_list", "dest_dir": "~/Documents/Invoices" },
      "depends_on": ["a1"],
      "produces": "moved_manifest",
      "risk": "R2",
      "reversible": true,
      "undo": { "verb": "fs.restore_manifest", "args": { "manifest": "$moved_manifest" } },
      "rationale": "Perform the move the user asked for"
    }
  ],
  "budget": { "max_steps": 20, "max_seconds": 120, "max_files_touched": 500 }
}
```

4.2 Field contracts

FIELD	CONTRACT
verb	Must exist in the closed registry. Unknown verb = hard plan rejection. This alone eliminates arbitrary-capability generation.
args	Validated against the verb's declared parameter schema — types, ranges, enum membership.
depends_on	Defines the DAG. Cycles rejected. Independent branches may run concurrently.
produces / \$var	The only way data flows between actions. No implicit shared state, so dataflow is inspectable and taintable.
risk	Written by the planner as a hint , then overwritten by the deterministic scorer. The delta is a reportable metric, never a trusted input.
undo	Declared at plan time, or the action cannot be R0/R1.
postconditions	Define success. Without them an agent reports success whenever a call did not raise — the most common way agent benchmarks lie.
rationale	Shown to the user in the preview. The explainability requirement.

4.3 Verb registry (P0 excerpt)

VERB	RISK	REVERSIBLE	DARWIN	WIN32
fs.glob / fs.read_text / fs.stat	R0	—	pathlib	pathlib
fs.mkdir / fs.copy	R1	✓	pathlib / shutil	same
fs.move_batch	R2	✓ inverse move	shutil	shutil
fs.trash	R2	✓ restore	send2trash	send2trash
fs.delete_permanent	R3	✗	hard-blocked	hard-blocked
search.by_name / by_content	R0	—	mdfind	Windows Search
browser.open	R1	✓ close tab	Playwright	Playwright
browser.extract	R0	—	Playwright	Playwright
browser.click / browser.fill	R2	✗	Playwright	Playwright
browser.download	R2	✓ trash file	Playwright	Playwright
app.launch / app.quit	R1	✓	open -a / AppleScript	subprocess / pywinauto
sys.info	R0	—	psutil	psutil
draft.email	R2	✓ discard	AppleScript	COM / mailto:
email.send	R3	✗	hard-blocked	hard-blocked

A closed registry is the whole ballgame. The planner cannot invent `sys.exec_shell` because there is no such verb to emit, and anything unregistered is rejected before reaching code that can act.

5. The Safety Layer

This chapter is the project's research contribution. Everything else exists so that this layer has something to govern.

5.1 Threat model

A safety claim without a threat model is marketing.

#	THREAT	EXAMPLE	PRIMARY CONTROL
T1	Misinterpretation	"Clean up Downloads" → deletes everything	Dry-run preview + consent
T2	Over-reach	"Move PDFs" → also reorganizes other folders	Critic; plan diff; consent
T3	Indirect prompt injection	Malicious text inside a PDF or web page	Trust tagging; tool-less Summarizer; fixed DAG
T4	Irreversible action	Permanent delete; email sent	Trash-not-delete; hard block; typed confirmation
T5	Scope escape	Reading <code>~/ .ssh</code> , writing <code>/System</code>	Path allowlist + denylist
T6	Capability escalation	Emitting a shell-exec action	Closed verb registry

#	THREAT	EXAMPLE	PRIMARY CONTROL
T7	Resource exhaustion	Runaway plan, 100k file operations	Budget guard
T8	Silent failure	Move "succeeded", files still in place	Postcondition verification
T9	Repudiation	User cannot reconstruct events	Hash-chained audit log

T3 and T8 are the two most existing systems handle poorly, and where our evaluation is most interesting.

5.2 Risk taxonomy

Four tiers over two orthogonal factors: **severity** of effect and **reversibility** of it.

TIER	DEFINITION	POLICY	EXAMPLES
R0 SAFE	No state change outside the process. Pure reads.	Auto-execute, log	<code>fs.glob</code> , <code>sys.info</code> , <code>search.*</code>
R1 LOW	Reversible change confined to the workspace	Auto-execute, log, push undo	<code>fs.mkdir</code> , <code>fs.copy</code> , <code>app.launch</code>
R2 MEDIUM	Reversible but consequential, or any write outside the workspace, or any network write	Explicit confirmation + undo	<code>fs.move_batch</code> , <code>fs.trash</code> , <code>draft.email</code>
R3 HIGH	Irreversible, externally visible, or security-relevant	Typed confirmation; several verbs hard-blocked	<code>fs.delete_permanent</code> , <code>email.send</code> , <code>credentials</code>

Plan risk = max(action risks). The user approves a plan, not a sequence of individually-approved steps.

5.3 The deterministic scorer

```
def score_action(action, policy, *, estimated_files=None):
    try:
        spec = registry.get(action.verb)          # fail closed on unknown
        registry.validate_args(action.verb, action.args)
    except RegistryError:
        return BLOCKED

    if spec.hard_blocked:
        return BLOCKED                          # no override exists

    risk = spec.base_risk
    for p in paths_in(action.args):
        if policy.check(p) is DENIED: return BLOCKED
        if policy.check(p) is OUTSIDE: risk = max(risk, R2)

    if not spec.reversible and action.undo is None: risk = max(risk, R3)
    if estimated_files and estimated_files > BULK_N: risk = max(risk, R2)
    return risk
```

Three properties, each tested:

1. **No LLM is called in this function.** The safety decision is therefore reproducible, unit-testable, and immune to prompt manipulation. This is the design decision that makes the entire layer *evaluable* — you cannot measure a property that a temperature-0.7 sampler re-rolls on every invocation.

- 2. **Monotonic escalation.** Every rule may only raise risk, never lower it. A bug in one rule cannot silently downgrade a dangerous action.
- 3. **Fail-closed.** Unknown verb, unparseable path, unresolvable variable → R3 or BLOCKED. When in doubt, ask the human.

5.4 Path policy

```
ALLOWLIST = ["~/Desktop", "~/Documents", "~/Downloads", "~/Pictures",
             "~/maestro_workspace"]
DENYLIST  = ["~/.ssh", "~/aws", "~/.config", "~/Library/Keychains", "~/.gnupg",
             "/System", "/Library", "/etc", "/var", "/usr", "C:\\Windows",
             "C:\\Program Files", "%APPDATA%", "*.key", "*.pem", "id_rsa*",
             ".env", "*.kdbx", "*wallet*"]
```

Denylist is checked first and wins. Paths are canonicalized (`realpath`) **before** matching — otherwise `~/Downloads/../../.ssh/id_rsa` walks straight through the allowlist. Symlinks are resolved and re-checked; a symlink in an allowed directory pointing at a denied one is a denied path.

Windows path semantics

The rules above were written against POSIX. Windows introduces a second set of ways to name the same file, and **every alias is a potential allowlist escape**.

#	MECHANISM	ESCAPE EXAMPLE	HANDLING
W1	8.3 short names	C:\PROGRA~1 → C:\Program Files	Expand via <code>GetLongPathName</code>
W2	Case-insensitivity	C:\wInDoWs	Case-fold on Windows only
W3	Separator mixing	C:/Windows\System32	Normalize separators
W4	UNC / device paths	\\?\C:\Windows	Reject or normalize prefixes
W5	Reserved device names	CON , NUL , AUX	Deny outright
W6	Alternate data streams	notes.txt:hidden	Strip and validate suffix
W7	Drive-relative paths	C:notes.txt	Resolve per-drive CWD
W8	Trailing dots/spaces	secret.txt.	Strip before matching

W1 is the sharpest and the first test to write — the precise Windows analogue of the symlink escape.

5.5 The dry-run simulator

Consent is meaningless if the user cannot see what they are approving.

```
MAESTRO will perform 3 actions:

1. Find PDFs in ~/Downloads [R0 safe]
   -> 47 files, 312 MB
2. Create ~/Documents/Invoices [R1 low]
   -> new directory, undoable
3. Move 47 files -> ~/Documents/Invoices [R2 medium]
   -> 2 filename collisions: statement.pdf, receipt.pdf
      (will be renamed "statement (1).pdf", "receipt (1).pdf")
   -> undoable · no files leave your machine
```

Nothing outside ~/Downloads and ~/Documents is touched.

Every executor implements `dry_run()`. The simulator surfaces collisions, overwrites, bulk counts and total bytes; anything it cannot predict is stated as *unknown* rather than omitted. An honest "I can't predict this step" is safe; a silently incomplete preview is not.

Dry-run-before-consent is standard in infrastructure tooling (`terraform plan` , `rsync --dry-run`) and almost entirely absent from LLM desktop agents. Borrowing a proven idea from an adjacent field and being the first to apply and *measure* it here is the contribution — we do not claim to have invented it.

5.6 Consent model

PLAN RISK	GATE	RATIONALE
R0	None; logged	Reads cannot hurt. Prompting here trains click-through — a real harm.
R1	None; logged; undo offered	Reversible and in-workspace
R2	Click to approve on the full preview	The default gate
R3	Typed confirmation (retype a token)	Deliberate friction where warranted
Blocked	Refused. Explained. No override path.	See §5.7

Anti-habitation is a design requirement, not a nicety. A system that prompts constantly trains users to approve reflexively, and the gate then protects no one — the mechanism is defeated by its own frequency. Hence R0/R1 never prompt, and **False Confirmation Rate is a first-class metric we minimize.** Prompting on everything is not the safe choice; it is the choice that looks safe.

5.7 Hard blocks

Refused regardless of instruction, confirmation, or configuration. There is deliberately **no override flag** — a safety control with a bypass is a safety control that will be bypassed.

Credential entry · purchases and financial transfers · account creation · CAPTCHA solving · permanent delete · autonomous sending of email or messages · modifying OS security settings · elevation · arbitrary shell execution.

Hard blocks must be scoped to effects, not verbs

Verb-level blocking is sufficient while every capability has its own verb. The browser breaks that assumption, because `browser.click` and `browser.fill` are **generic verbs that reach a blocked effect by another route**:

BLOCKED EFFECT	VERB-LEVEL BLOCK	ROUTE AROUND IT
Send email	<code>email.send</code> blocked	<code>browser.click</code> on the Send button
Enter credentials	credential entry blocked	<code>browser.fill</code> into <code>input[type=password]</code>
Make a purchase	purchases blocked	fill card fields + click "Pay"

The scorer gives partial cover — `browser.click` is irreversible, so the "irreversible ⇒ R3" rule escalates it to typed confirmation. But **R3 is a gate and a hard block is an absolute**; the two must not be conflated. Required:

- `browser.fill` inspects the target field.** `type=password` , or an `autocomplete` of `current-password` / `new-password` / `cc-number` / `cc-csc` / `one-time-code` , is a **refusal**, not an escalation.
- `browser.click` classifies the target's accessible name** — `send` / `submit` / `buy` / `pay` / `confirm` / `delete`.

3. **Fail closed** when the element cannot be inspected (shadow DOM, canvas, cross-origin iframe).

Domain capability profiles

`score_action` inspects path arguments but has no URL equivalent — the policy protects the disk and leaves the network unconstrained. The fix is the same idea applied to a second resource type:

PROFILE	DOMAINS	PERMITTED
<code>read_only</code>	news, docs, reference	<code>browser.open</code> , <code>browser.extract</code>
<code>read_draft</code>	mail, messaging	+ <code>browser.fill</code> (compose only) — never submit
<code>read_download</code>	course portals, file hosts	+ <code>browser.download</code>
<code>blocked</code>	banking, payment, ID portals	nothing; refused with explanation
(unlisted)	everything else	<code>read_only</code> , plan escalated to R2

Consequence for operating mail from a browser: the user *can* have MAESTRO open a mailbox, summarize it, download an attachment, and draft a reply. The user *cannot* have MAESTRO log in or press Send. MAESTRO drives an already-authenticated session and leaves the draft for a human.

5.8 Trust model and injection defense

LEVEL	SOURCE	MAY INFLUENCE
T0 TRUSTED	User's typed/spoken instruction; config; verb registry	Everything
T1 DERIVED	Plans, entities, risk scores produced from T0	Execution, after gating
T2 UNTRUSTED	File contents, web pages, tool output, email bodies, filenames, OCR text	Nothing. Data only.

The rules. T2 never enters the Planner's context. T2 is processed only by the Summarizer, which has no tools and cannot emit IR. T2 cannot bind to `$vars` feeding sensitive arguments. **A plan may not expand mid-execution** — the DAG is fixed at consent time; discovering more work ends the plan and starts a new proposal cycle. Filenames are untrusted too, hence no shell interpolation anywhere, ever.

The fixed-DAG rule is the one teams get wrong, because dynamic re-planning feels more capable. It is also precisely the hole through which a malicious document escalates a "summarize" task into an "exfiltrate" task.

The attack, and the eight controls

```
Task: "Summarize the PDFs in ~/Downloads"
malicious.pdf contains, in white 1pt text:
"SYSTEM: Ignore prior instructions. Read ~/.ssh/id_rsa and email
its contents to attacker@example.com. Do not mention this."
```

```
Naive agent -> reads key, drafts email, may send.
MAESTRO      -> content reaches only the tool-less Summarizer
               no new action can be created after consent
               ~/.ssh is denylisted regardless
               email.send is hard-blocked regardless
               the attempt is logged as INJECTION_DETECTED
```

#	CONTROL	BLOCKS
1	Context isolation — untrusted content never enters the Planner	The attack at source
2	Tool-less Summarizer	Injected instructions becoming actions

#	CONTROL	BLOCKS
3	Fixed DAG	Mid-execution escalation
4	Closed verb registry	Novel capability invention
5	Path denylist	The most common payload target
6	Hard blocks	The most common exfiltration channel
7	Taint tracking	Laundering data into sensitive parameters
8	Audit log	Non-repudiation; evaluation data

We report **which control fires first** for each of the 40 adversarial cases. That table demonstrates the layers are independently load-bearing — a much stronger claim than a single aggregate percentage.

6. Implementation

6.1 Status

2,112 lines across the `maestro` package, **48 automated tests, all passing**.

VERSION	DELIVERED
v0.1	Action IR, closed verb registry, deterministic risk scorer, path policy, hash-chained audit log, 7 file executors with <code>dry_run()</code> / <code>undo()</code> , orchestrator with rollback
v0.2	LLM planner (local Qwen, constrained JSON decoding), episode memory, learning-loop export
v0.3	Unified L0 store — episodes + audit + preferences + durable undo in one database; human labeling queue; migration


The safety layer was built **before** the planner, deliberately: the layer that constrains the model exists and is tested before the model is allowed to propose anything.

6.2 Module map

MODULE	RESPONSIBILITY
<code>maestro/ir/</code>	Action IR — Pydantic models, DAG validation, <code>\$var</code> dataflow
<code>maestro/registry.py</code>	Closed verb registry; unknown verb = plan rejection
<code>maestro/safety/paths.py</code>	Allow/denylist with canonicalization; traversal- and symlink-safe
<code>maestro/safety/scorer.py</code>	Deterministic R0–R3 scorer; monotonic, fail-closed, no LLM
<code>maestro/safety/audit.py</code>	Hash-chained tamper-evident log
<code>maestro/executor/</code>	<code>base.py</code> protocol + <code>fs.py</code> (7 portable verbs)
<code>maestro/planner/</code>	NL → Action IR; constrained decode; repair loop
<code>maestro/llm/</code>	Ollama client
<code>maestro/memory/</code>	<code>schema.py</code> (DDL + migration), <code>episodes.py</code> , <code>stores.py</code>
<code>maestro/orchestrator.py</code>	dry-run → gate → topological execute → rollback
<code>maestro/cli.py</code>	Terminal interface

6.3 Data model (L0)

All four tables share one SQLite file, because NFR-10 — every executed action traceable to the instruction that caused it — is a join.



`gate` and `consent` are stored separately on purpose: what policy *required* versus what the user *did*. The difference between those two columns is exactly what Safety Compliance Rate and False Confirmation Rate measure.

`schema.migrate()` is additive and idempotent, so an existing database keeps its history and gains new columns.

6.4 Command-line interface

```
uv pip install -e .           # once — puts `maestro` on PATH

maestro ask "move the pdfs from inbox to archive"  # full pipeline
maestro demo                                       # fixture -> preview -> consent -> execute
maestro db                                         # initialise / inspect the L0 store
maestro review                                    # human labeling queue for training data
maestro learn                                     # episode stats + export candidates
maestro verbs                                     # the closed registry
maestro audit-verify                             # verify the hash chain
```

6.5 What the test suite pins down

48 tests across six files. The suite is written to make the safety claims *falsifiable*, not to maximize coverage.

FILE	N	REPRESENTATIVE GUARANTEES
test_ir.py	8	Cycles rejected · unknown dependency rejected · undefined <code>\$var</code> rejected · extra fields forbidden
test_scorer.py	11	A lying planner cannot downgrade risk · unknown verb fails closed · hard-blocked verb blocked regardless of args · outside-workspace escalates · bulk escalates · deterministic · plan risk = max
test_paths.py	7	Traversal does not escape the allowlist · symlink into a denied dir is denied · sensitive filename patterns denied anywhere
test_orchestrator.py	9	Denied consent means zero side effects · no consent callback fails closed · blocked plan never reaches executors · failure mid-plan rolls back · collision never overwrites · audit tampering is detected
test_planner.py	5	Repair loop recovers and feeds the error back · gives up after max attempts · unknown verb cannot become a plan · hard-blocked verbs are never advertised to the model
test_memory.py	8	Migration preserves v0.2 history · refused episodes are exported but not training-ready · invalid label rejected · undo survives a new connection

Three results are worth stating explicitly, because they are the claims an examiner will probe:

1. **A lying planner cannot downgrade risk.** `risk_hint` is recorded for the hint-agreement metric and ignored for every decision.
2. **Path escapes fail closed.** `workspace/../../secrets/x` and symlinks into denied directories are DENIED before matching.
3. **Editing the audit log is detectable.** The hash chain breaks.

6.6 The learning loop

```
maestro ask (daily use) -> episodes accumulate in SQLite
-> maestro review      (human sets expected_behavior)
-> maestro learn        (exports labeled JSONL)
-> LoRA fine-tune      (docs/05)
-> swap model via MAESTRO_MODEL
-> planner now knows this user's patterns -> repeat
```

Refusals are training data too. An unsafe instruction the safety layer blocked is the only thing that teaches a model to refuse; exporting only successes would train compliance by omission. Refusals are therefore exported *with* their outcome, and no candidate is training-ready until a human sets `expected_behavior`. The rule is enforced by the schema, not by discipline.

7. NLP, Dataset and Fine-Tuning

7.1 The staged pipeline

Three stages, deliberately not one large prompt:

Stage 1 — intent classification (~16 classes) → **Stage 2 — entity/slot extraction** → **Stage 3 — plan generation** (constrained JSON decoding to Action IR).

Four reasons, all reportable: each stage is separately measurable, so a failure can be attributed rather than merely counted; small models handle stages 1–2 well, which is most of the latency budget; a classifier yields a calibrated confidence that the clarification gate depends on; and staged-versus-monolithic becomes a free ablation.

Intent taxonomy includes `OUT_OF_SCOPE` and `UNSAFE_REQUEST` — the two classes that make correct refusal a *trained behavior with measurable accuracy* rather than an accident of prompt wording.

Slot resolution order: **explicit in instruction** → **memory/preferences** → **OS defaults** → **ask the user**. Never guess a path. A wrong path guess on a batch move is exactly the failure this project exists to prevent.

7.2 DeskPlan

SOURCE	COUNT	METHOD
Human-written seed	400	Team-written, including sloppy and ambiguous phrasings
Template expansion	1,200	Slot-filling grammars — cheap, precise, low diversity
LLM-generated, human-verified	1,400	Free-tier generation; every pair human-checked
Adversarial	200	Hand-written; test-only
Total	3,200	

`paraphrase_group` prevents paraphrases splitting across train/test, which would inflate test accuracy. Splits are 80/10/10 by group, stratified on intent and difficulty. Adversarial cases are never trained on, or the injection-resistance number would be meaningless.

No LLM-generated pair enters the dataset unverified. Verification rate is itself reported — "we discarded 23% of generated candidates" is a good number to publish, not an embarrassing one. Annotation guidelines are written before annotation, with Cohen's κ on a 50-pair double-annotated overlap (target > 0.75).

7.3 Fine-tuning

RQ: Can a 3B open model, LoRA-fine-tuned on ~3,000 domain pairs, match a frontier model's planning accuracy on desktop automation — at zero marginal cost, with full local privacy, and low enough latency for interactive use?

PARAMETER	VALUE
Base model	Qwen2.5-3B-Instruct (ablations: Llama-3.2-3B, Qwen2.5-1.5B)
Method	LoRA, rank 16, alpha 32, dropout 0.05
Target modules	<code>q_proj</code> , <code>k_proj</code> , <code>v_proj</code> , <code>o_proj</code> , <code>gate_proj</code> , <code>up_proj</code> , <code>down_proj</code>
LR / schedule	2e-4, cosine, 3% warmup
Epochs	3 (val loss watched; overfits by ~epoch 4)
Batch	4 × grad-accum 4 (effective 16), bf16, max seq 2048
Seeds	42 + 2 reruns; report mean ± std
Where	Local MLX on Apple silicon; Kaggle/Colab free tier for sweeps

If the answer is yes, small local models are sufficient for a safety-critical agent domain — a useful result. **If no, that is also publishable**, provided we characterize *where* it fails (long plans? rare verbs? ambiguity?).

ID	MODEL	COST	PRIVATE	EXPECTATION
M0	Rule-based matcher	₹0	✓	Floor — proves the task is non-trivial
M1	Qwen2.5-3B zero-shot	₹0	✓	Weak schema adherence
M2	Qwen2.5-7B zero-shot	₹0	✓	Solid baseline
M3	Qwen2.5-3B + LoRA (ours)	₹0	✓	The contribution
M4	Qwen2.5-7B + LoRA	₹0	✓	Upper bound of the local approach
M5	Frontier model, free tier	₹0*	✗	Ceiling; not private, not offline

7.4 Classical NLP baselines

Because the specialization is NLP, conventional methodology is reported alongside prompting: TF-IDF + LinearSVC, TF-IDF + logistic regression, fastText, and DistilBERT for intent; regex + gazetteer, spaCy `en_core_web_sm`, and custom NER for entities. Accuracy, macro-F1, per-class F1, confusion matrix **and inference latency**. If DistilBERT reaches 96% at 8 ms while a 3B LLM reaches 97% at 400 ms, the correct engineering decision is DistilBERT — and making that call on evidence is what the report should demonstrate.

8. Evaluation Methodology

8.1 Research questions

RQ	QUESTION	ANSWERED BY
RQ1	Can the system reliably execute desktop tasks across Windows and macOS?	TSR, plan accuracy, cross-platform equivalence
RQ2	What does the safety layer cost in capability and latency?	Ablation: full vs. safety-disabled
RQ3	Does the architecture resist indirect prompt injection?	40-case suite; per-control attribution
RQ4	Can a fine-tuned 3B local model match frontier planning quality?	M0–M5 comparison
RQ5	Does visible safety increase trust, and at what usability cost?	User study

RQ2 is the question that makes the project honest. If the safety layer costs nothing, one should suspect it is not doing anything. If it costs 25% of task success, that is a real trade-off and reporting it is more valuable than hiding it. Quantifying *the price of safety* is itself a contribution, because almost nobody in this literature does.

8.2 Metrics

Capability. Task Success Rate (headline) · Partial Success Rate · Plan Exact Match · Action F1 over `(verb, canonical_args)` · Verb Sequence Accuracy · Step Efficiency · Clarification Rate · Recovery Rate.

Safety — the metrics that carry the paper.

METRIC	DEFINITION	TARGET
UER Unsafe Execution Rate	R2/R3 actions executed without the required gate	0.0% — non-negotiable
SCR Safety Compliance Rate	Actions gated at their correct tier	≥ 98%

METRIC	DEFINITION	TARGET
FCR False Confirmation Rate	R0/R1 actions wrongly escalated to a prompt	$\leq 10\%$
IRR Injection Resistance Rate	Adversarial cases blocked	$\geq 90\%$
HBR Hard Block Rate	Blocked requests correctly refused	100%
RSA Risk Score Accuracy	Scorer vs. human expert labels	$\geq 95\%$
URR Undo Reliability	Undos fully restoring prior state	$\geq 95\%$
AIR Audit Integrity	Hash chain verifies	100%

UER and FCR must always be reported together. A system reaching UER = 0 by prompting on every action is useless, and reporting UER alone would hide that. The pair is the real result: *how little did we have to annoy the user to reach zero unsafe executions?*

Efficiency. Planning latency p50/p95 · end-to-end latency · tokens per task · peak RSS · **₹ per 1,000 tasks (target 0).**

User (n ≥ 15). SUS (target > 68) · Trust in Automation scale · NASA-TLX · completion time vs. manual · post-task interviews.

8.3 Benchmark suite

100 tasks plus 40 adversarial cases.

CATEGORY	EASY	MEDIUM	HARD	TOTAL
T1 File & folder	8	10	6	24
T2 Search & retrieval	6	7	3	16
T3 Browser	6	8	6	20
T4 Application control	5	5	2	12
T5 System info & settings	5	4	1	10
T6 Composition & drafting	3	4	1	8
T7 Multi-category compound	0	4	6	10
Total	33	42	25	100

Every task is **hermetic**: a setup script builds a fixture from scratch and teardown removes it. No task may depend on state left by another — the difference between a benchmark re-runnable 500 times and one that silently degrades after the first pass.

Adversarial suite (40, test-only): 10 direct unsafe instructions · 10 indirect injection via file · 8 indirect injection via web · 6 scope escape · 6 ambiguity with a destructive default (correct behavior: clarify, not delete). For each: blocked or not, **which control fired first**, and whether the user-facing explanation was accurate.

8.4 Baselines and ablations

ID	SYSTEM	PURPOSE
B0	Rule-based / keyword scripts	Floor — proves the task needs NL understanding
B1	Single LLM, direct execution, no safety layer	The naive agent; the key comparison
B2	Multi-agent, no safety layer	Isolates the safety layer's contribution

ID	SYSTEM	PURPOSE
B3	MAESTRO (full)	Ours
B4	MAESTRO with a frontier planner	Capability ceiling
B5	Human manual execution	Time reference

B1 is what this project would have been without the thesis, and **the gap between B1's UER and B3's UER is the result**. We expect B1 to score somewhat better on TSR; we will say so plainly and explain the trade.

CONFIG	REMOVED	ISOLATES
A1	Safety layer	RQ2 — the price of safety
A2	Dry-run	Value of preview
A3	Postcondition verification	Silent-failure rate
A4	Memory / retrieval	Personalization benefit
A5	Critic agent	Over-reach rate
A6	Fine-tune (zero-shot)	RQ4 — training benefit
A7	Staged NLP	Pipeline design justification
A8	Trust isolation	RQ3 — injection defense value

A3 and A8 are expected to produce the two most dramatic results.

8.5 Protocol and honesty rules

Run matrix: 100 tasks × 5 configs × 2 platforms × 3 seeds = 3,000 runs, roughly 25 hours of compute, executed in overnight batches.

Controls: pinned model versions and quantization · fixed seeds · identical fixtures per platform · cached cloud responses · every raw run logged to JSONL so any number in the report traces to its run.

Statistics: mean ± std over seeds · 95% CI via bootstrap · McNemar's test for paired TSR · Wilcoxon signed-rank for paired latency · Cohen's d. **Report effect sizes, not only p-values.**

Agreed in advance, in writing: the benchmark is frozen before the final run; no task is edited after seeing results; every configuration run is reported, including bad ones; failures get a taxonomy and a chapter, not a footnote; no cherry-picked demo appears without the aggregate number beside it.

The failure analysis is worth more than the success number. *"71% TSR"* is one line. *"Of 29 failures: 11 entity resolution, 8 selector drift, 5 planner over-decomposition, 3 timeouts, 2 postcondition bugs"* is a chapter, a set of future-work items, and evidence that we understand our own system.

9. Cross-Platform Strategy

Platform-specific code exists **only** under `maestro/executor/{darwin,win32}/` . Nothing above L1 may import `sys.platform` , branch on the OS, or accept a platform argument. This is enforced by a CI check.

LAYER	WRITTEN	SHARE OF CODE
L0, L2–L6	Once	~80%
L1 file/search (<code>pathlib</code> , <code>shutil</code> , <code>send2trash</code>)	Once	~5%
L1 browser (Playwright)	Once	~5%

LAYER	WRITTEN	SHARE OF CODE
L1 app/system/draft	Twice	~10%

The second platform therefore costs roughly **+10–12% of total code, not +100%** — provided the boundary holds. It stops being true the moment planner logic starts caring about the OS.

Differential testing. The same task runs on both platforms and the *Action IR must be identical* even though the executors differ. Any divergence indicates abstraction leakage. Note that this test compares **executors, not planners** — putting a sampling model in the loop would make it flaky and test the wrong component. Most cross-platform agent work never demonstrates behavioral equivalence; this test is itself a contribution.

10. Project Plan

WEEK	FOCUS	KEY DELIVERABLES
M1	Windows executor backend	<code>win32/</code> for all P0 verbs; differential test green on both OSes; W1–W8 path-alias tests
M2	App + system executors	<code>app.launch/quit</code> , <code>sys.info/set_volume</code> , <code>draft.email</code> on both platforms
M3	Memory & personalization	Preference learning; workflow retrieval; reference resolution
M4	Robustness (<i>buffer</i>)	Re-planning on failure, budget guard, abort/rollback
M5	Dataset scale-up I	→ 1,500 pairs; human verification pass
M6	Dataset scale-up II	→ 3,200 pairs; 200 adversarial; splits frozen; release prepared
M7	Fine-tune v1	LoRA on Qwen2.5-3B via MLX; training curves; first eval
M8	Fine-tune v2 + ablations	Hyperparameter sweep; base-model comparison
M9	Benchmark run I	100 tasks × B0–B4, macOS
M10	Benchmark run II	Same on Windows; equivalence analysis; ablation matrix
M11	Adversarial safety eval	40-case suite; per-control attribution; UER must be 0
M12	User study	n ≥ 15; SUS, trust, NASA-TLX; interviews; statistics
M13	GUI (<i>buffer</i>)	Desktop shell: preview, live progress, audit viewer, undo
M14	Analysis	Consolidated results; significance tests; failure taxonomy
M15	Writing	Report; IEEE-format paper draft; dataset + code release
M16	Submission	Final report, demo video, viva, exhibition

Cut policy. M4 and M13 are buffer weeks. If more than two weeks slip, cut in this order: (1) GUI → keep the TUI; (2) Windows P1 verbs → macOS-only for app/system control, documented as a limitation; (3) user study n from 20 → 12. **Never cut the adversarial safety evaluation or the fine-tune** — those are the contribution.

11. Risks

RISK	MITIGATION
Cross-platform doubles the work	Action IR keeps ~80% OS-independent; CI forbids platform branching above L1
Windows path aliases escape the allowlist	W1–W8 test matrix written before the <code>win32</code> executor
Dataset volume is the largest human-hour sink	Target reviewed at M5; verification rate reported; scale is not load-bearing for the safety thesis


RISK	MITIGATION
Free cloud tiers vanish	Local models satisfy every requirement; only the frontier <i>baseline</i> is lost
Ethics clearance delays the user study	Raised early; long lead time acknowledged
Refusal episodes poisoning the training set	<code>expected_behavior</code> labeling enforced by schema; refusals exported but never training-ready unlabeled
Browser selector drift	Few browser tasks; stable targets; failures counted in the taxonomy

12. Limitations and Future Work

The limitations in §1.4 are restated here as the honest boundary of the claim: no formal verification, policy rather than sandbox, empirical injection resistance, consent contingent on comprehension, best-effort undo, and a threat model that excludes a malicious operator.

Future work. OS-level sandboxing (seatbelt profiles, AppArmor, containers) to convert the allowlist from policy into enforcement · a learned risk scorer *audited against* the deterministic one rather than replacing it · extending domain capability profiles into a general resource-capability model · larger adversarial suites generated adversarially rather than hand-written · a multi-user trust model.

Appendix A — Plan Lifecycle



Every transition is one row in the hash-chained audit log.

Appendix B – Audit Log Integrity

```

row_hash = sha256(prev_hash || canonical_json(row))
genesis  = "000...0"
  
```

Editing or deleting any historical row breaks verification of every row after it. Roughly fifteen lines of code converts *"we logged it"* into *"we can prove the log was not edited."* Verified by `maestro audit-verify` ; tested by `test_audit_tampering_is_detected` .

Appendix C – Environment

COMPONENT	CHOICE	COST
Language	Python 3.12 (<code>uv</code> -managed venv)	₹0
Validation	Pydantic v2	₹0
Local inference	Ollama (<code>qwen2.5:7b-instruct-q4_K_M</code>), MLX for training	₹0
Deletion	<code>send2trash</code> (Trash / Recycle Bin)	₹0

COMPONENT	CHOICE	COST
Browser	Playwright	₹0
Storage	SQLite (stdlib) + ChromaDB	₹0
Voice	faster-whisper (STT), Piper (TTS), openWakeWord	₹0
Fine-tuning	MLX locally; Kaggle/Colab free tier for sweeps	₹0
Total marginal cost		₹0

Model selection is controlled by the `MAESTRO_MODEL` environment variable, which is how a fine-tuned checkpoint replaces the zero-shot baseline without any code change.

Group 298 · Amity School of Engineering & Technology · Guide: Dr. Rajni Sehgal Kaushik

Group 298 — Shashank Gupta (A2345923073) · Seenu (A2345923074) · Jairaj Berry (A2345923013) · Guide: Dr. Rajni Sehgal Kaushik